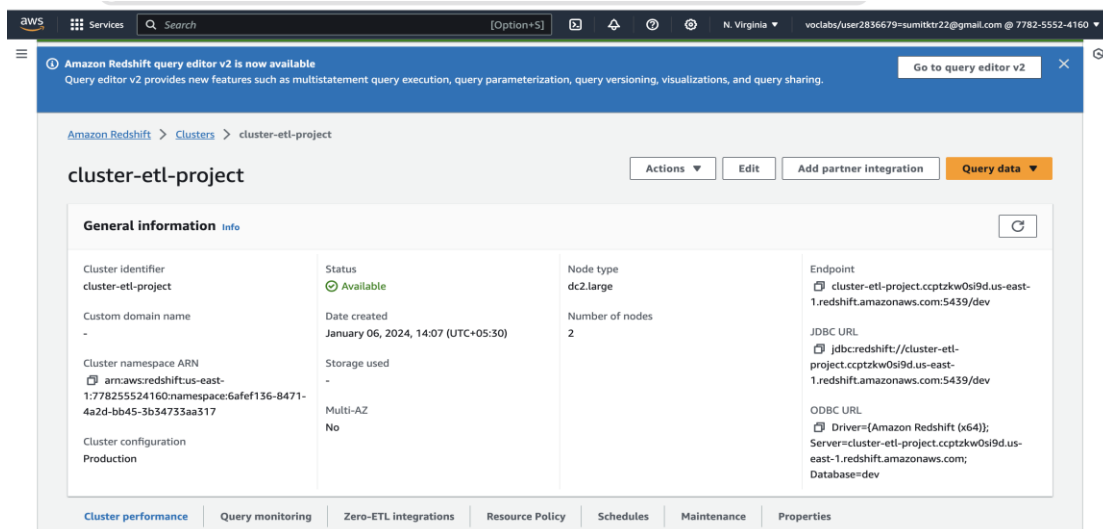
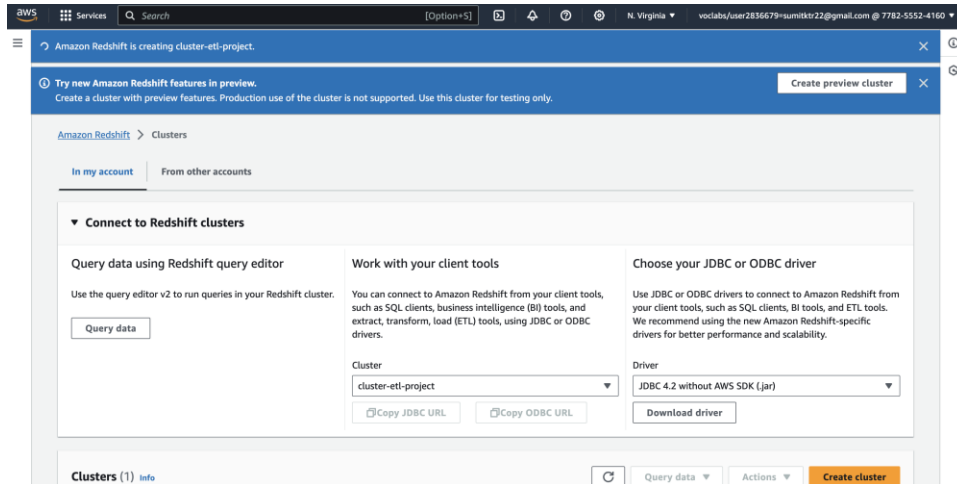


# Redshift cluster creation

Screenshots depicting the configuration of the Redshift cluster I've established:



Services

Search

[Option+S]

N. Virginia

voclabs/user2836679@sumitkr22@gmail.com @ 7782-5552-4160

Amazon Redshift query editor v2 is now available

Query editor v2 provides new features such as multistatement query execution, query parameterization, query versioning, visualizations, and query sharing.

Go to query editor v2

Amazon Redshift > Clusters > cluster-etl-project

Actions Edit Add partner integration Query data

cluster-etl-project

General information

Cluster identifier cluster-etl-project	Status Available	Node type dc2.large	Endpoint cluster-etl-project.cptkzw0s9d.us-east-1.redshift.amazonaws.com:5439/dev
Custom domain name -	Date created January 06, 2024, 14:07 (UTC+05:30)	Number of nodes 2	JDBC URL jdbc:redshift://cluster-etl-project.cptkzw0s9d.us-east-1.redshift.amazonaws.com:5439/dev
Cluster namespace ARN arn:aws:redshift:us-east-1:778255524160:namespace:6afef136-8471-4a2d-bb45-3b54733aa317	Storage used -		ODBC URL Driver=[Amazon Redshift (v64)]; Server=cluster-etl-project.cptkzw0s9d.us-east-1.redshift.amazonaws.com; Database=dev
Cluster configuration Production	Multi-AZ No		

Cluster performance Query monitoring Zero-ETL integrations Resource Policy Schedules Maintenance Properties

Services

Search

[Option+S]

N. Virginia

voclabs/user2836679@sumitkr22@gmail.com @ 7782-5552-4160

Production

east-1.redshift.amazonaws.com; Database=dev

Cluster performance Query monitoring Zero-ETL integrations Resource Policy Schedules Maintenance Properties

Database configurations

Edit admin credentials Rotate encryption keys Edit

Database name dev	Parameter group Defines database parameter and query queues for all the databases. default.redshift-1.0	Encryption Disabled	Audit logging Disabled
Port 5439	SSH ingestion setting (cluster public key) ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQCE 5CCzUff+TkhV3Sm61eq9K2Yb+L9di2Gul3I ncQIE+XV9F4z+qWmLg5Od1a3RlBmjyEH 5zwVvHpuVH/OKTvwQsGFRZQamFNuF1a saYBoTbk+QT7QB8CwkoUSQBLDfRWWnf eU5vOCoFgVveSMYahzLIHIOvK8dWSAXW aIG/G7m9OROPtp1W6EXbgJFKJ9rlGncpbEX hwyUkJOSEqKLhB+dxdbV8Q+KAbdUglCU1c TP8AP0e4scu64u8uXIMXKZ2Zfzm+YMQq Q7p3RkXUlgQX4/wtkdwwfzpxqxJf6yD5M Vl/qeoyvE2fj+cz8gNAeVoMmkun1aH Amazon-Redshift	AWS KMS key ID -	
Admin user name awsuser			
Admin password ARN arn:aws:secretsmanager:us-east-1:778255524160:secret:redshiftcluster-etl-project-awsuser-q7eHGn			

Services

Search

[Option+S]

N. Virginia

voclabs/user2836679@sumitkr22@gmail.com @ 7782-5552-4160

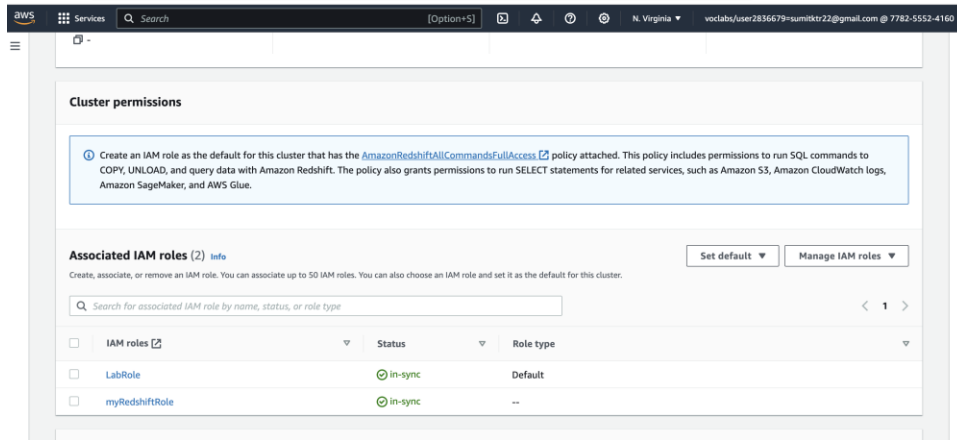
Network and security settings

Edit

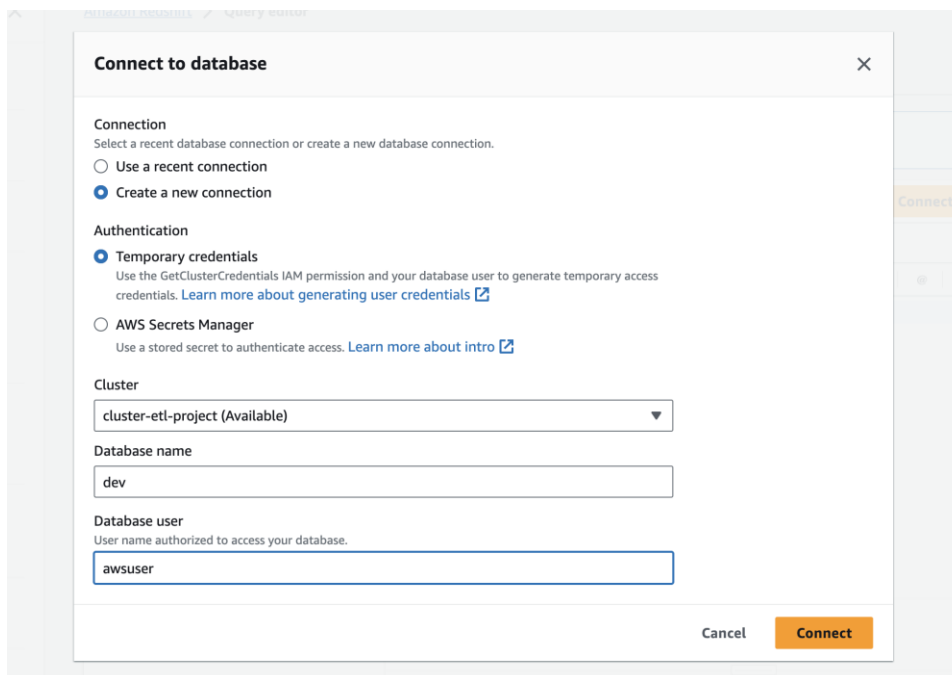
Virtual private cloud (VPC) vpc-0a62c993e28b2b981	Availability Zone us-east-1f	VPC security group Specify which instances and devices can connect to the cluster. sg-0c9b387f753de6dc8	Publicly accessible Allow connections from outside the VPC. Disabled
Subnet group default	Enhanced VPC routing Disabled	IP address type -	
Endpoint URL -			

Cluster permissions

Create an IAM role as the default for this cluster that has the AmazonRedshiftFullAccess policy attached. This policy includes permissions to run SQL commands to COPY, UNLOAD, and query data with Amazon Redshift. The policy also grants permissions to run SELECT statements for related services, such as Amazon S3, Amazon CloudWatch logs, Amazon SageMaker, and AWS Glue.



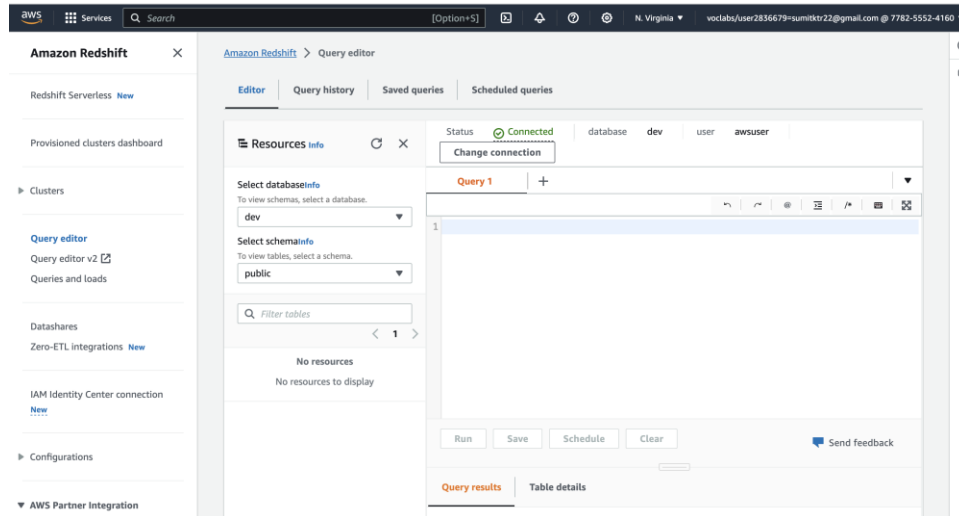
## Connection to the database :



Establishing a database within the Redshift cluster and executing queries to create both dimension and fact tables.

Query for creating schema:

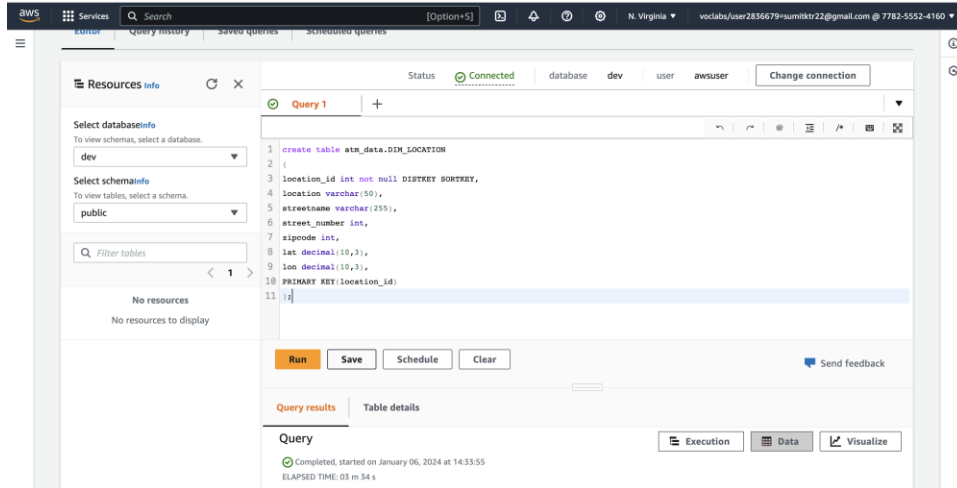
**create schema atm\_data;**



Constructing queries for the creation of diverse dimension and fact tables, ensuring proper assignment of primary and foreign keys.

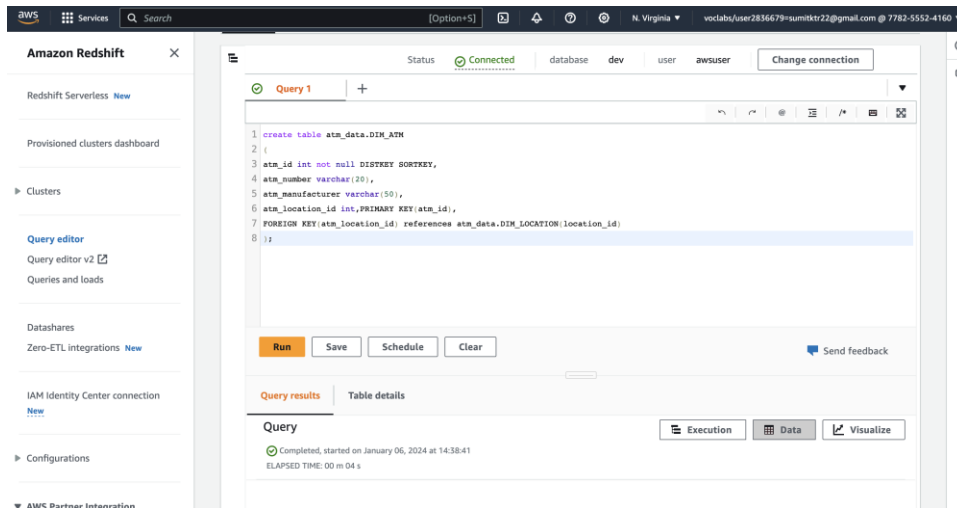
Creating location dimension table

```
create table atm_data.DIM_LOCATION
(  
  location_id int not null DISTKEY SORTKEY,  
  location varchar(50),  
  streetname varchar(255),  
  street_number int,  
  zipcode int,  
  lat decimal(10,3),  
  lon decimal(10,3),  
  PRIMARY KEY(location_id)
```



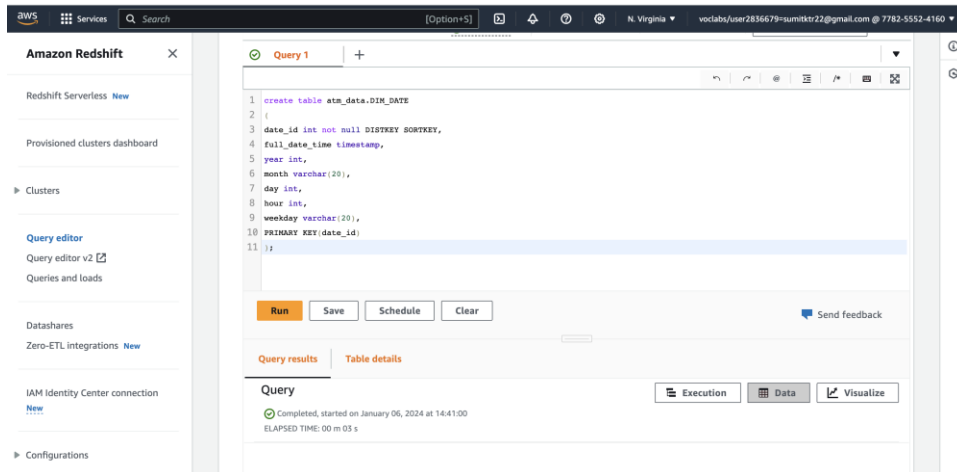
## Creating atm dimension table

```
create table atm_data.DIM_ATM
(
  atm_id int not null DISTKEY SORTKEY,
  atm_number varchar(20),
  atm_manufacturer varchar(50),
  atm_location_id int,
  PRIMARY KEY(atm_id),
  FOREIGN KEY(atm_location_id) references atm_data.DIM_LOCATION(location_id)
);
```



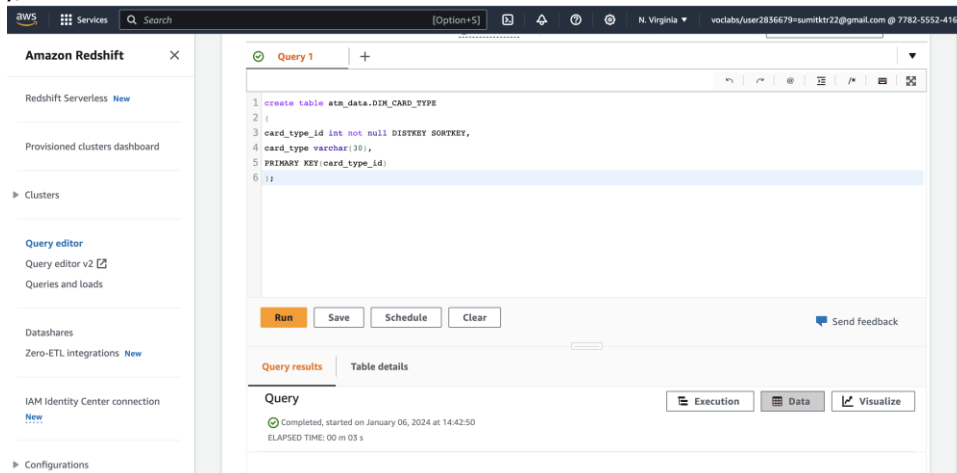
## Creating date dimension table

```
create table atm_data.DIM_DATE
(
  date_id int not null DISTKEY SORTKEY,
  full_date_time timestamp,
  year int,
  month varchar(20),
  day int,
  hour int,
  weekday varchar(20),
  PRIMARY KEY(date_id)
);
```



## Creating card type dimension table

```
create table atm_data.DIM_CARD_TYPE
(
  card_type_id int not null DISTKEY SORTKEY,
  card_type varchar(30),
  PRIMARY KEY(card_type_id)
);
```



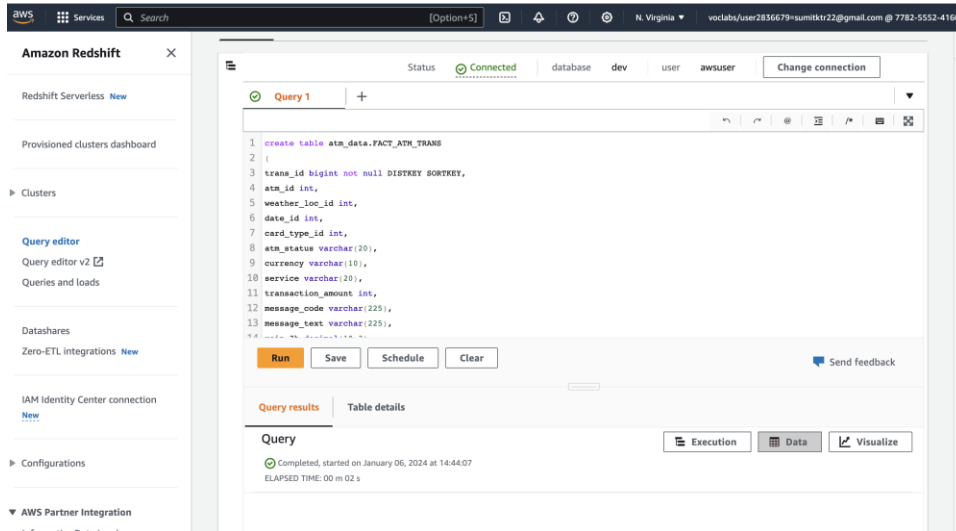
## Creating atm transactions fact table

```
create table atm_data.FACT_ATM_TRANS
(
  trans_id bigint not null DISTKEY SORTKEY,
  atm_id int,
  weather_loc_id int,
  date_id int,
  card_type_id int,
  atm_status varchar(20),
  currency varchar(10),
  service varchar(20),
  transaction_amount int,
  message_code varchar(225),
  message_text varchar(225),
  rain_3h decimal(10,3),
  clouds_all int,
  weather_id int,
  weather_main varchar(50),
  weather_description varchar(255),
  PRIMARY KEY(trans_id),
  FOREIGN KEY(weather_loc_id) references atm_data.DIM_LOCATION(location_id),
  FOREIGN KEY(atm_id) references atm_data.DIM_ATM(atm_id),
  FOREIGN KEY(date_id) references atm_data.DIM_DATE(date_id),
  FOREIGN KEY(card_type_id) references atm_data.DIM_CARD_TYPE(card_type_id)
);
```

## Loading data into a Redshift cluster from Amazon S3 bucket

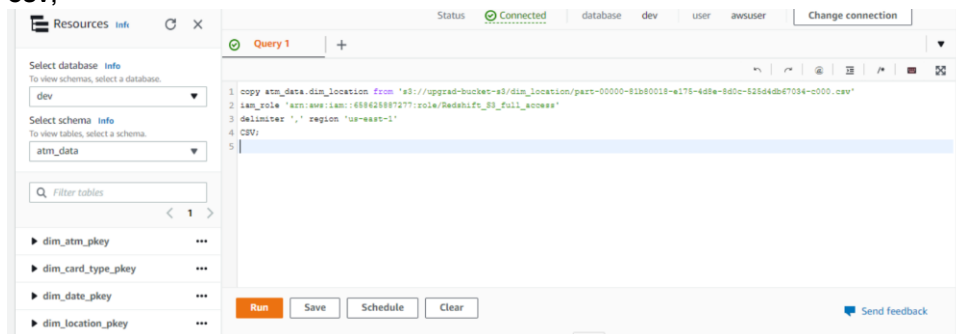
Queries to copy the data from S3 buckets to the Redshift cluster in the appropriate tables

- Copying the data to dim\_location table



## Copying the data to dim\_atm table

copy atm\_data.dim\_location from 's3://upgrad-bucket-s3/dim\_location/part-00000-81b80018-e175-4d8e-8d0c-525d4db67034-c000.csv'  
iam\_role 'arn:aws:iam::658625887277:role/Redshift\_S3\_full\_access'  
delimiter ',' region 'us-east-1'  
CSV;



## Copying the data to dim\_date table

copy atm\_data.dim\_atm from 's3://upgrad-bucket-s3/dim\_atm/part-00000-86ed9a65-6e1e-

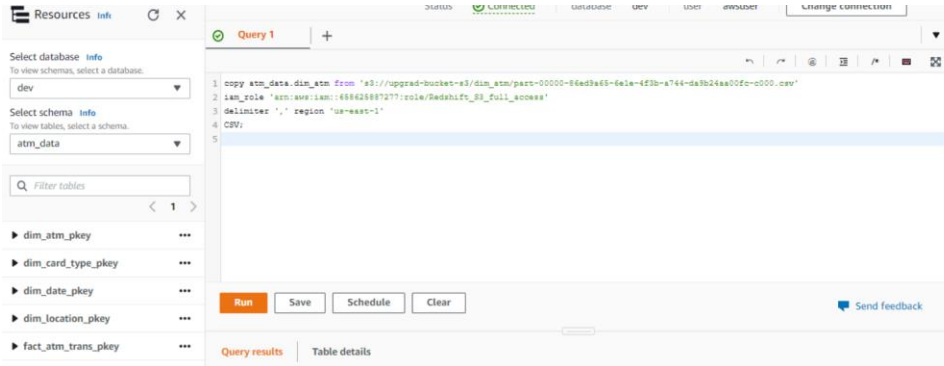


**4f3b-a744-da9b24aa00fc-c000.csv'**

**iam\_role 'arn:aws:iam::658625887277:role/Redshift\_S3\_full\_access'**

**delimiter ',' region 'us-east-1'**

**CSV;**



**copy atm\_data.dim\_date from 's3://upgrad-bucket-s3/dim\_date/part-00000-beb40dab-**

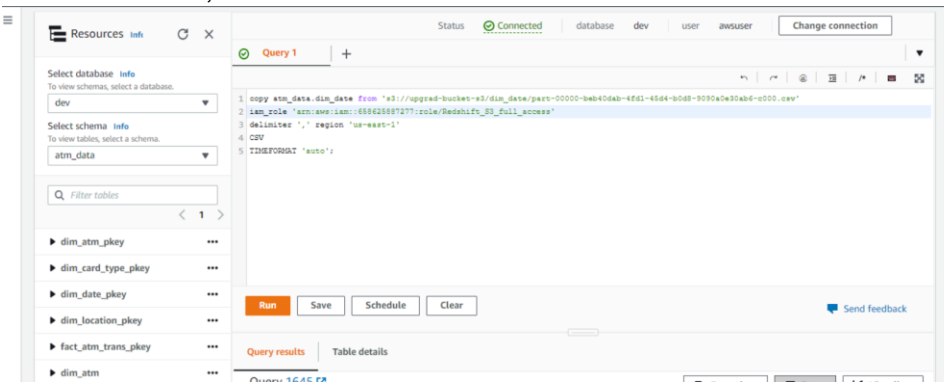
**4fd1-45d4-b0d8-9090a0e30ab6-c000.csv'**

**iam\_role 'arn:aws:iam::658625887277:role/Redshift\_S3\_full\_access'**

**delimiter ',' region 'us-east-1'**

**CSV**

**TIMEFORMAT 'auto';**



## Copying the data to dim\_card\_type table

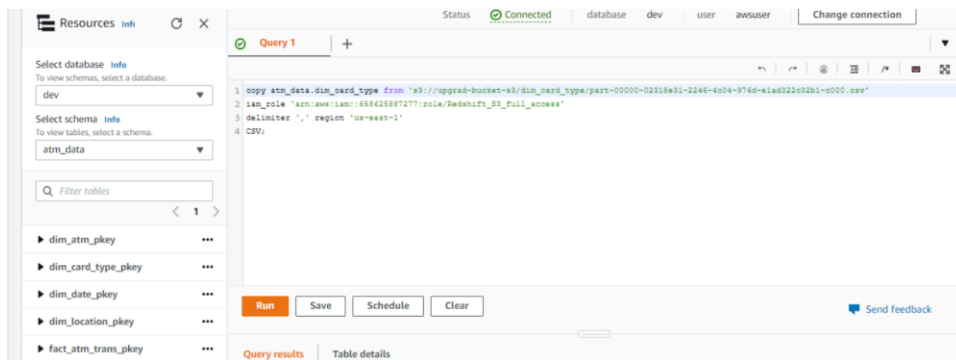
**copy atm\_data.dim\_card\_type from 's3://upgrad-bucket-s3/dim\_card\_type/part-00000-**

**02318e31-2246-4c04-976d-e1ad322c32b1-c000.csv'**

**iam\_role 'arn:aws:iam::658625887277:role/Redshift\_S3\_full\_access'**

**delimiter ',' region 'us-east-1'**

**CSV;**



## Copying the data to fact\_atm\_trans table

**copy atm\_data.fact\_atm\_trans from 's3://upgrad-bucket-s3/fact\_atm\_trans/part-00000-fb72e768-f589-4caa-bd26-1bc6fcbbdb8a-c000.csv'**  
**iam\_role 'arn:aws:iam::658625887277:role/Redshift\_S3\_full\_access'**  
**delimiter ',' region 'us-east-1'**  
**CSV;**

